



Lessons from defending the indefensible


Marek Majkowski

marek@cloudflare.com @majek04

Denial of service (DoS)



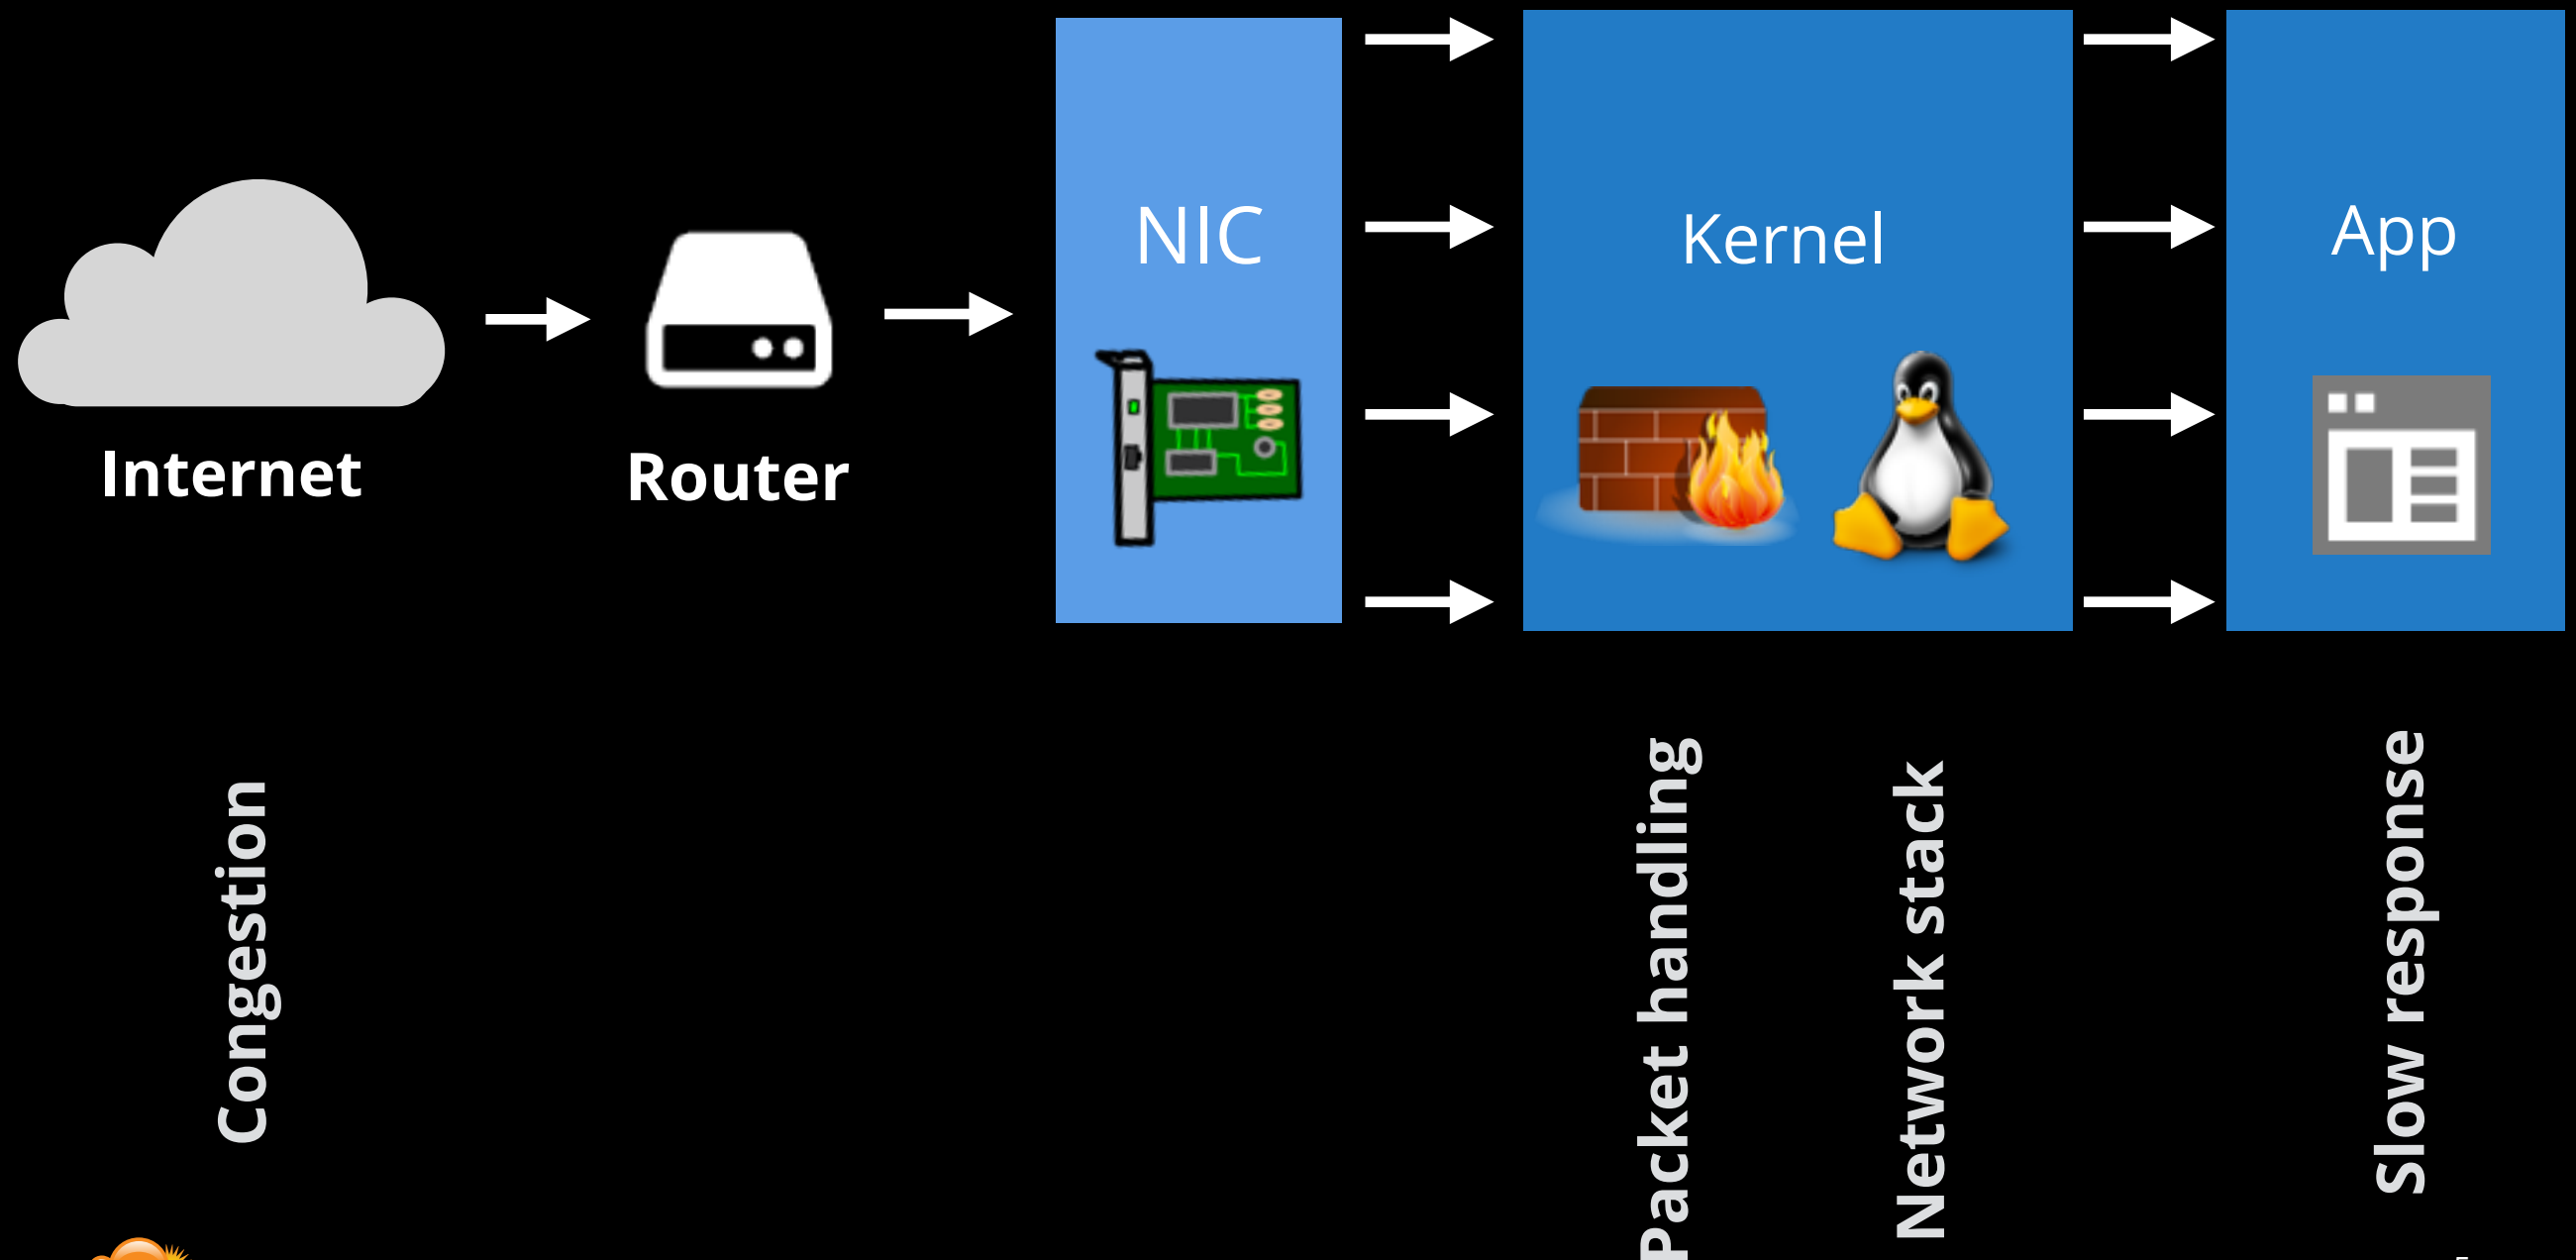
DoS is a hard problem




Unique view



Attack surface



Packets per second



Your Linux needs your help



∞ pps

Network congestion

Congestion

```
$ netstat -s
```

```
Tcp:
```


```
...
```

```
2291681363 segments send out
```

```
43887463 segments retransmitted
```

```
...
```

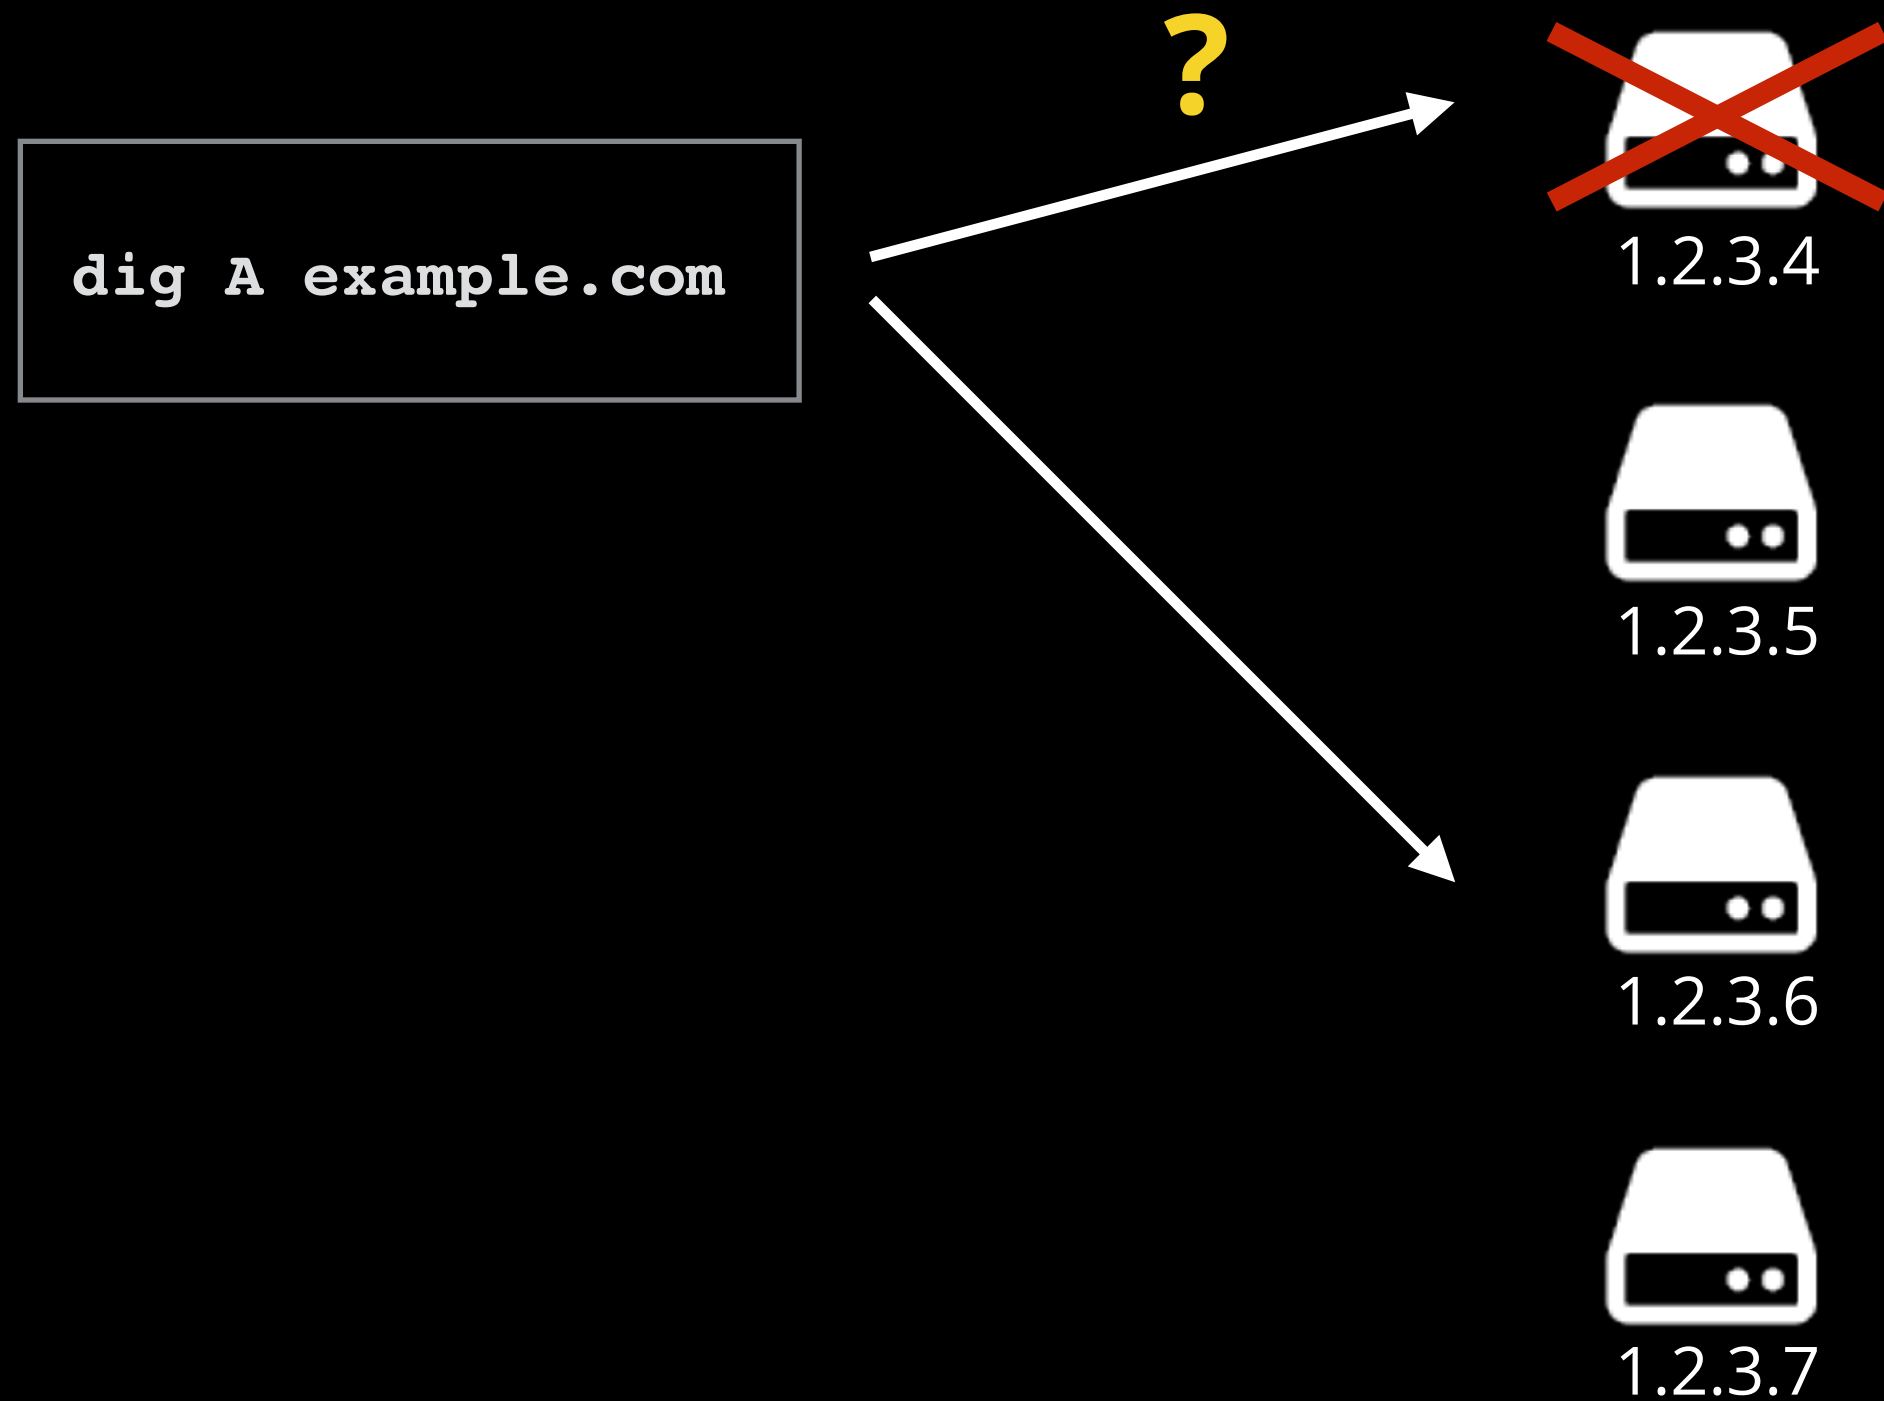
Congestion



BGP null routing

```
route 198.41.222.x/32 {  
    discard;  
    community [ 13335:666 13335:668 13335:36006];  
}
```

Application integration




High volume packet floods

Let it flow



Looks like this

Packets per second





Spoofed?



Drop!

UDP flood

```
IP 23.243.202.207.1076 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 98.151.75.108.32856 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 23.118.154.219.33894 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 23.242.35.159.1036 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 23.240.170.79.33842 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 23.241.212.223.2052 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 187.204.126.111.59011 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 24.24.164.88.2759 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 23.242.122.1.32778 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
IP 23.240.26.33.1043 > 1.2.3.4:53: 18516 updated% [b2&3=0x5450] [11825a]
```

Packet characteristics

- packet length
- source IP's
- source port
- IPID field
- payload



~1.2M pps

Matching on payload in
iptables



Payload matching with BPF

```
iptables -A INPUT \  
  --dst 1.2.3.4 \  
  -p udp --dport 53 \  
    -m bpf --bytecode "14,0 0 0 20,177 0 0 0,12 0 0  
0,7 0 0 0,64 0 0 0,21 0 7 124090465,64 0 0 4,21 0 5  
1836084325,64 0 0 8,21 0 3 56848237,80 0 0 12,21 0 1  
0,6 0 0 1,6 0 0 0" \  
  -j DROP
```



BPF bytecode

```
    ldx 4*([14]&0xf)
    ld #34
    add x
    tax
lb_0:
    ldb [x + 0]
    add x
    add #1
    tax
    ld [x + 0]
    jneq #0x07657861, lb_1
    ld [x + 4]
    jneq #0x6d706c65, lb_1
    ld [x + 8]
    jneq #0x03636f6d, lb_1
    ldb [x + 12]
    jneq #0x00, lb_1
    ret #1
lb_1:
    ret #0
```



Tcpdump expressions


- Originally: `tcpdump -n "udp and port 53"`
- xt_bpf implemented in 2013 by Willem de Bruijn
- Need to deal with BPF byte code
- Tools around it are scarce (tcpdump expressions)
- Tcpdump expressions are limited - for example matching valid DNS packets case insensitive
- Need to hand-craft BPF



BPF tools


- Open source:
 - <https://github.com/cloudflare/bpftools>
- Can match various DNS patterns:
 - `*.example.com`
 - `???.example.com`
 - `*{1-4}.example.com`
 - `--case-insensitive *.example.com`
 - `--invalid-dns`

Total dropped

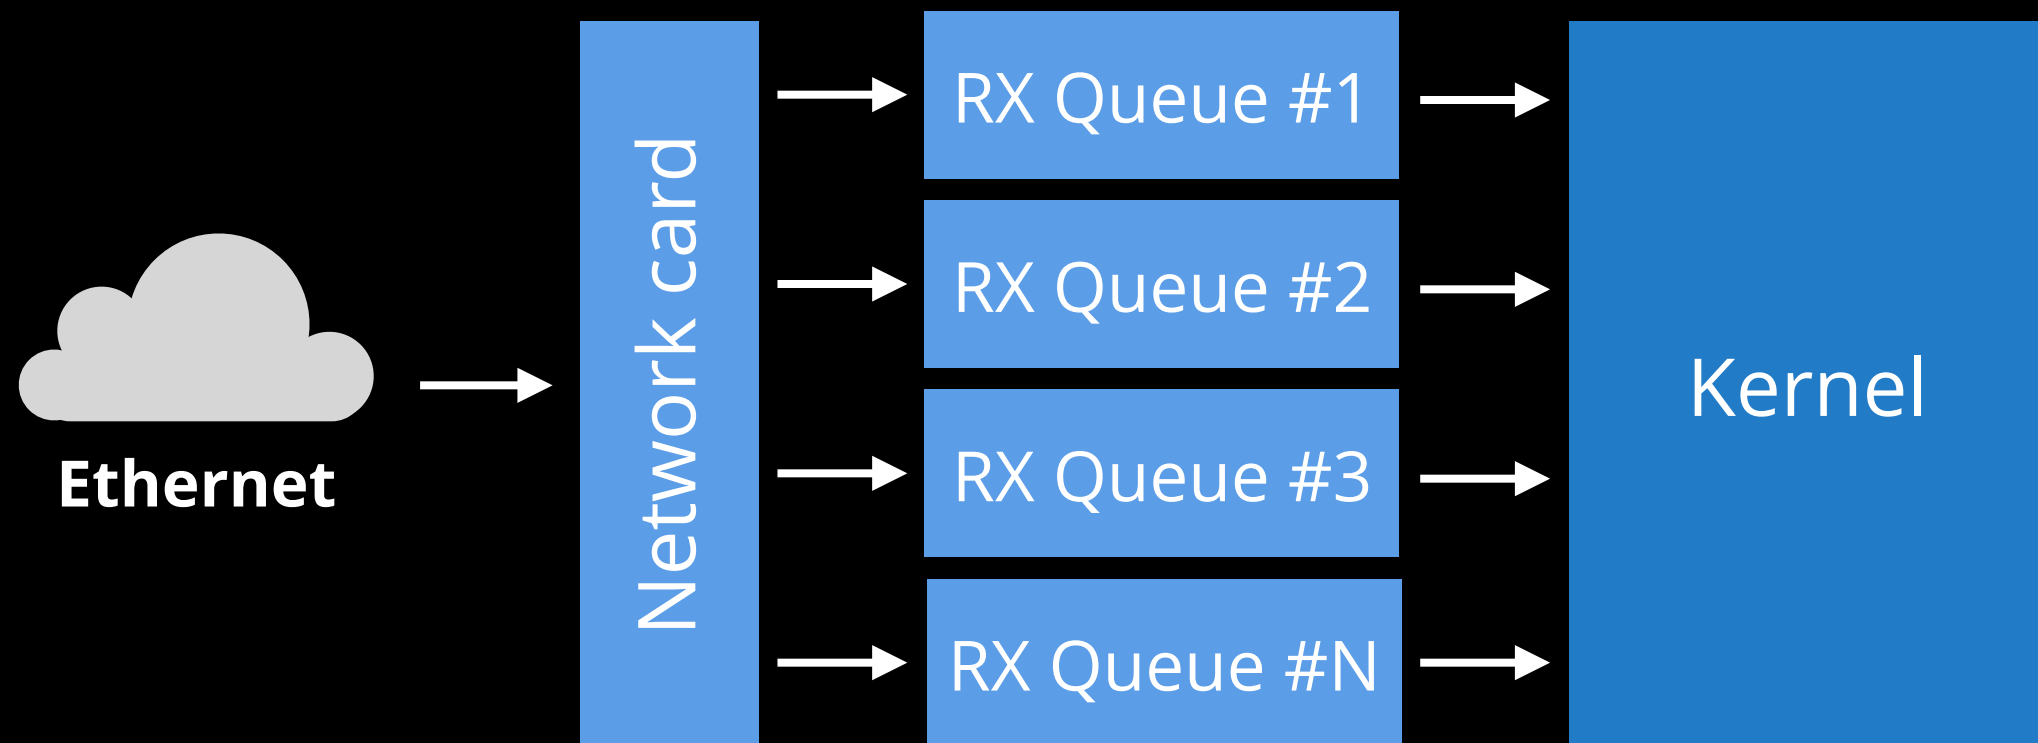



~4M pps

Payload matching close to NIC

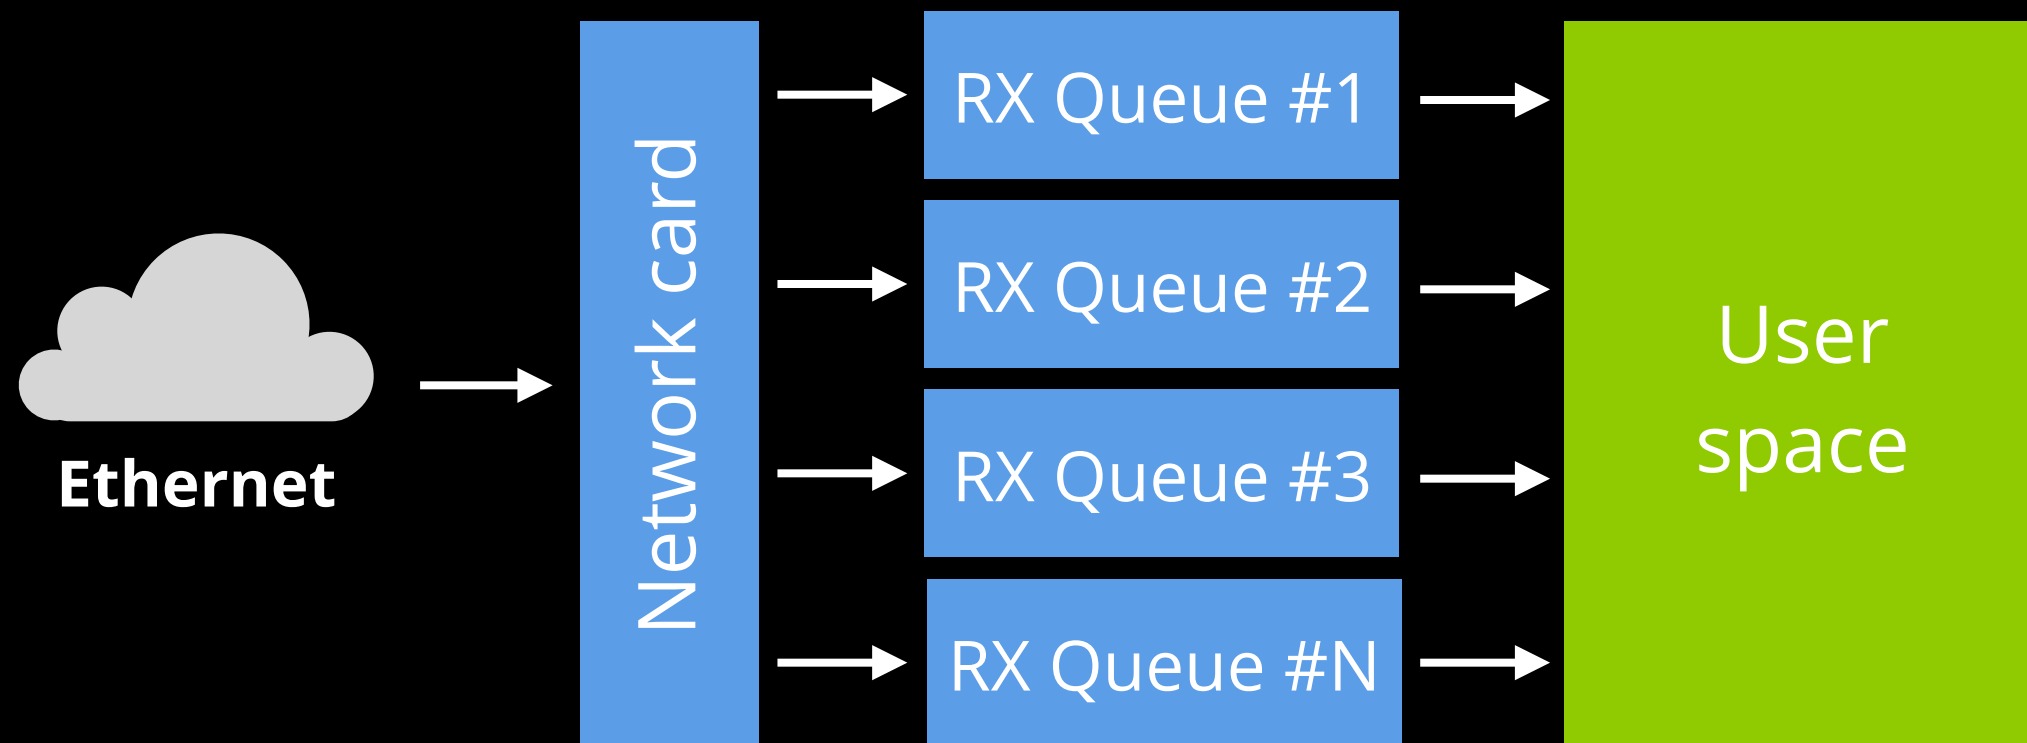



Modern NIC's



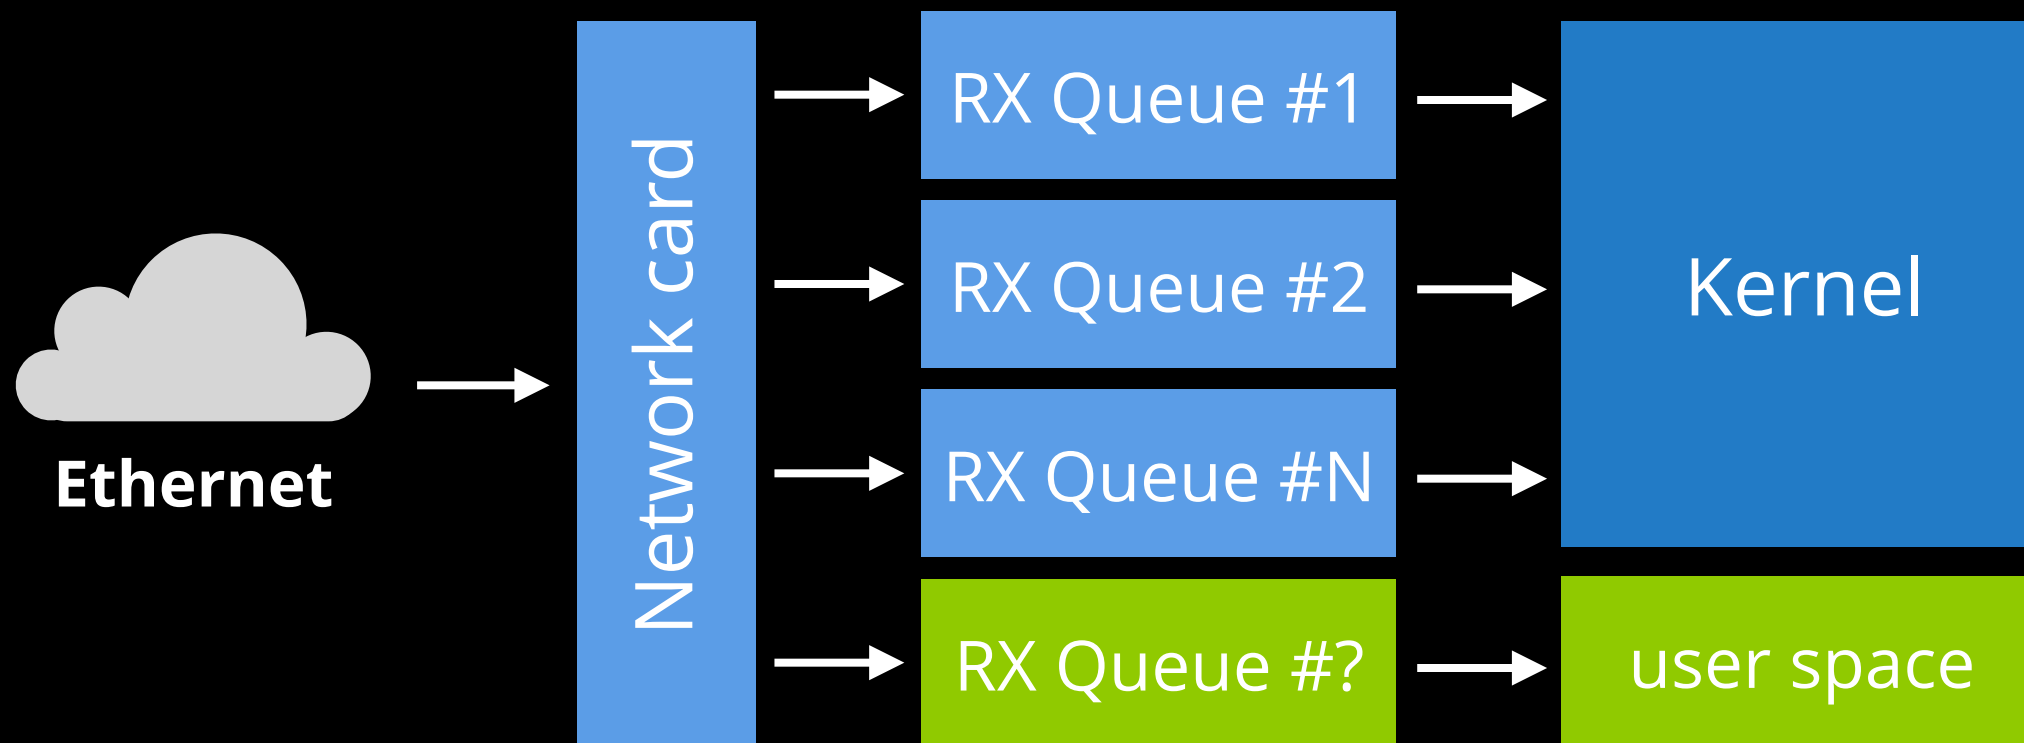



Traditional kernel bypass






Partial kernel bypas



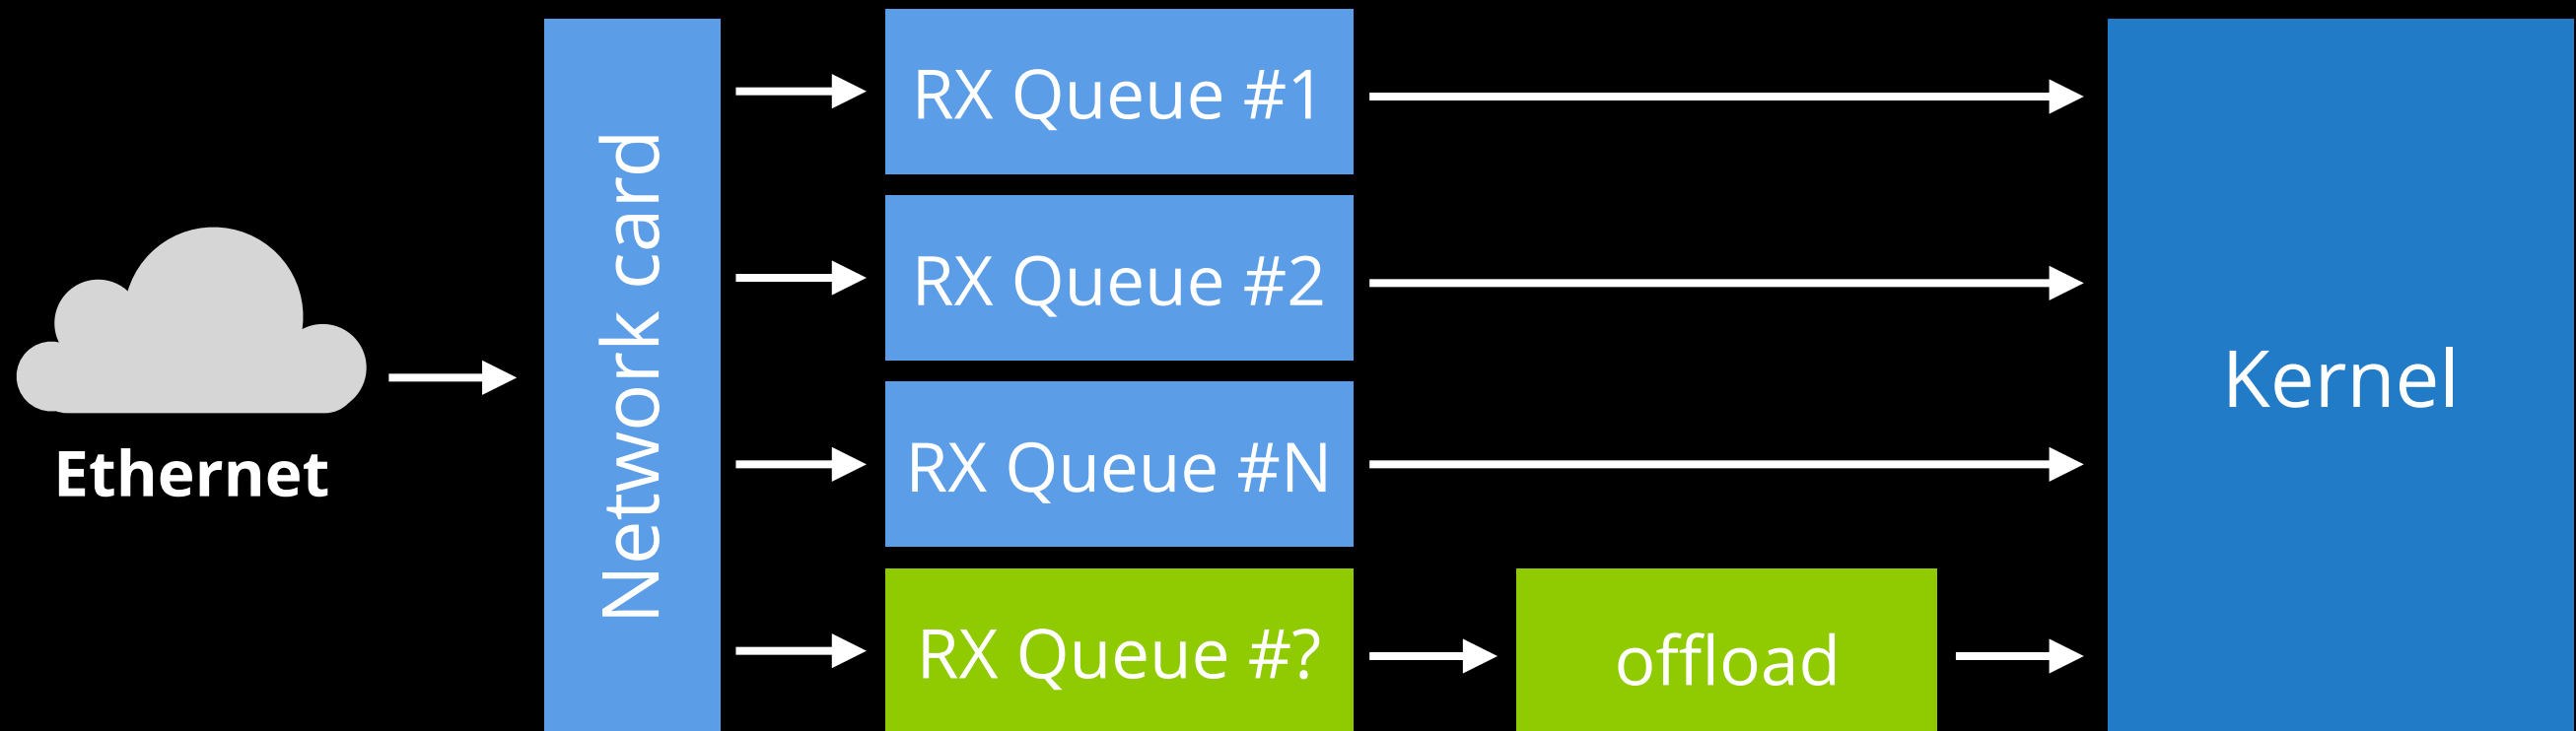


Partial kernel bypass

- Or EFVI for SolarFlares:
 - <http://www.openonload.org/>
- Open sourced netmap patch, tested on Intel:
 - <https://github.com/luigirizzo/netmap/pull/87>



Iptables offload



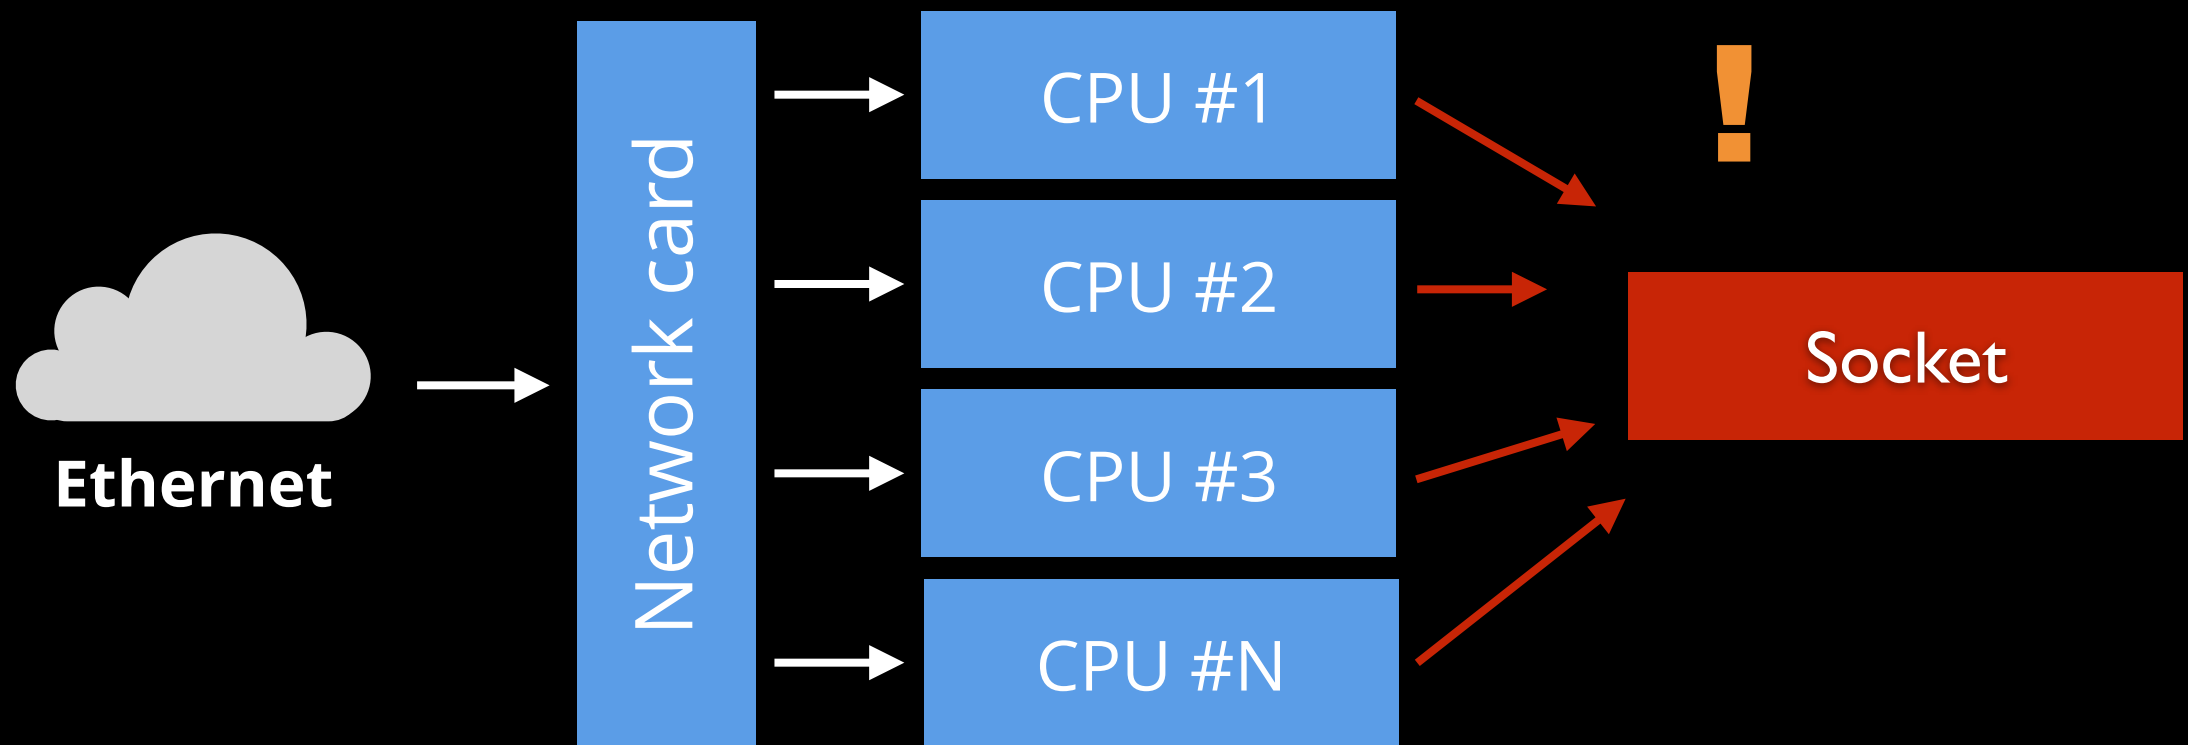
No characteristics:
Attacks against TCP/IP
network stack

ACK floods

```
IP 48.60.32.50.15244 > 1.2.3.4.80: Flags [P.], ack 1754729313, win 16153
IP 31.102.214.103.13396 > 1.2.3.4.80: Flags [P.], ack 1569851274, win 15707
IP 112.36.216.55.56515 > 1.2.3.4.80: Flags [P.], ack 2051477187, win 16102
IP 65.130.63.30.10341 > 1.2.3.4.80: Flags [P.], ack 2108282782, win 16112
IP 16.18.205.115.15962 > 1.2.3.4.80: Flags [P.], ack 1359019408, win 16119
IP 128.177.247.54.13752 > 1.2.3.4.80: Flags [P.], ack 1416531343, win 16102
IP 204.59.118.78.61528 > 1.2.3.4.80: Flags [P.], ack 348671255, win 16101
IP 119.195.142.20.3344 > 1.2.3.4.80: Flags [P.], ack 1917538144, win 16161
IP 70.197.6.24.39340 > 1.2.3.4.80: Flags [P.], ack 1920842431, win 16124
```

~0.3M pps

Fight for the lock





Statefull firewall - conntrack

```
iptables -A INPUT \  
  --dst 1.2.3.4 \  
  -m conntrack --ctstate INVALID \  
  -j DROP
```

```
sysctl -w net/netfilter/nf_conntrack_tcp_loose=0
```

~1.2M pps

Effective

- Works well against:
 - ACK
 - FIN
 - RST
 - X-mas
- What about SYN floods?


SYN floods

```
IP 94.242.250.109.47330 > 1.2.3.4:80: Flags [S], seq 1444613291, win 63243
IP 188.138.1.240.61454 > 1.2.3.4:80: Flags [S], seq 1995637287, win 60551
IP 207.244.90.205.17572 > 1.2.3.4:80: Flags [S], seq 1523683071, win 61607
IP 94.242.250.224.65127 > 1.2.3.4:80: Flags [S], seq 928944042, win 61778
IP 207.244.90.205.43074 > 1.2.3.4:80: Flags [S], seq 137074667, win 63891
IP 64.22.81.44.23865 > 1.2.3.4:80: Flags [S], seq 838596928, win 63808,
IP 188.138.1.137.23373 > 1.2.3.4:80: Flags [S], seq 593106072, win 60272
IP 207.244.90.205.39653 > 1.2.3.4:80: Flags [S], seq 47289666, win 63210
IP 208.66.78.204.64197 > 1.2.3.4:80: Flags [S], seq 1850809890, win 62714
IP 207.244.90.205.33108 > 1.2.3.4:80: Flags [S], seq 319707959, win 63351
IP 207.244.90.205.6937 > 1.2.3.4:80: Flags [S], seq 1591500126, win 63902
IP 213.152.180.151.60560 > 1.2.3.4:80: Flags [S], seq 1902119375, win 62511
IP 64.22.79.127.11061 > 1.2.3.4:80: Flags [S], seq 1456438676, win 62148
```

0M pps



SYN in Linux





SYN backlog size

1. Listen backlog size

```
listen(int sockfd, int backlog)
```

2. Listen backlog size capped by

```
sysctl -w net.core.somaxconn = 65535
```

3. SYN backlog capped with

```
sysctl -w net.ipv4.tcp_max_syn_backlog = 65535
```

4. Rounded to the *NEXT* power of two

127 --> 128 128 --> 256

SYN backlog churn

```
sysctl -w net.ipv4.tcp_synack_retries=1
```



SYN backlog overflow



- Normal case - DROP
- Fixed with SYN cookies

TCP: Possible SYN flooding on port 80

SYN cookies



sequence number:

5 bits $t \bmod 32$	3 bits MSS	24 bits $\text{hash}(\text{ip}, \text{port}, t)$
------------------------	---------------	---

Tip: TCP timestamps

```
sysctl -w net.ipv4.tcp_timestamps=1
```

timestamp:

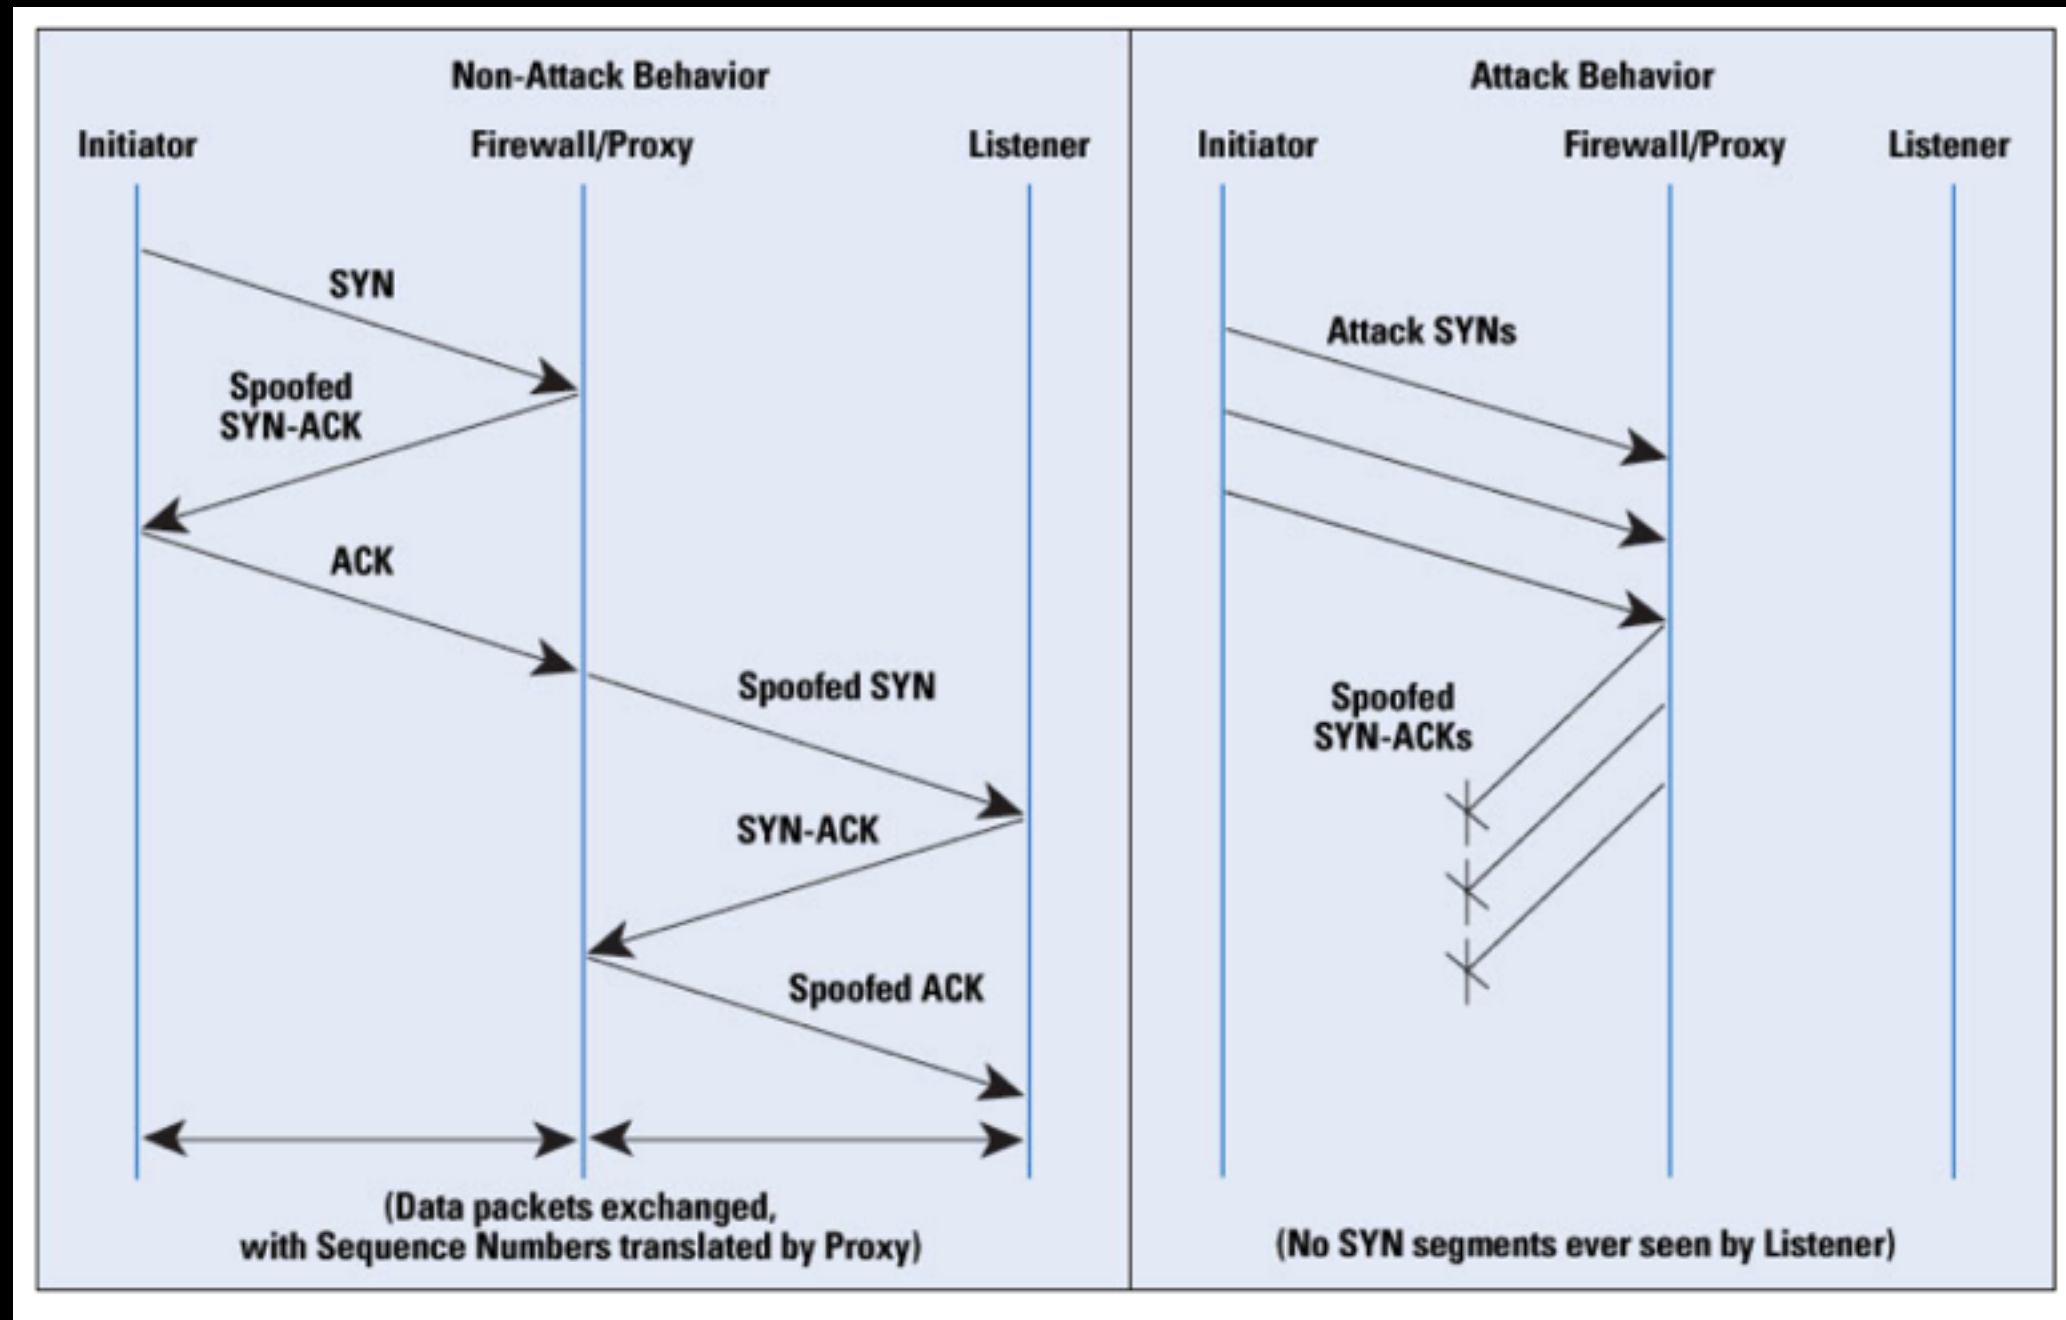
26 bits timestamp	1 bit ECN	1 bit SACK	4 bits wscale
----------------------	--------------	---------------	------------------

sequence number:

5 bits $t \bmod 32$	3 bits MSS	24 bits $\text{hash}(\text{ip}, \text{port}, t)$
------------------------	---------------	---

0.3M pps

Fight for the lock



(source: [Jesper Brouer presentation](#))

Recent changes

- The idea is to remove the LISTEN lock
 - Heavy refactoring of the SYN queue
- Submitted by Eric Dumazet in early October 2015
- Merged to net-next, will land in 4.4


Real connections from a
botnet

Real TCP/IP connections



Looks like this

Packets per second



Symptoms

- Connection count grows
- "Orphaned" sockets count grow
- "Time waits" growing

```
sysctl -w net.ipv4.tcp_max_orphans=262144  
sysctl -w net.ipv4.tcp_orphan_retries=1  
  
sysctl -w net.ipv4.tcp_max_tw_buckets=360000  
sysctl -w net.ipv4.tcp_tw_reuse=1  
sysctl -w net.ipv4.tcp_fin_timeout=5
```

IP reputation





Iptables to the rescue

- Ipset
 - Supports subnets
 - Manual blacklisting
- Hashlimits
 - Rate limit packets per subnet
 - Automatic blacklisting with timeout
- Connlimit
- TARPIT




Note on large botnets

- Make it a SYN flood
 - Disable HTTP keep-alives
- Blacklist IP's based on payload
 - Typo in request
 - BPF or string module for match + ipsets auto expiry


Note on large botnets

```
GET /forum.php HTTP/1.1
Accept: */*
Accept-Language: zh-cn
Accept-Encoding: gzip, deflate
User-Agent: Mozilla/5.0 (compatible; Baiduspider/2.0;...
Host: www.example.com:80
Connection: Keep-Alive
```

310k RPS, 650k uniques



Tip: sflow



Centralized Sflow

```
$ tailsflow -i sflow | tcpdump -n -r - -c 10 'vlan and ip'
reading from file -, link-type EN10MB (Ethernet)
IP 10.11.8.17.8070 > 10.11.8.82.24982:
IP 10.16.8.95.8070 > 10.16.10.139.33176: 18:55:22.345369
IP 70.215.131.237.3232 > 104.16.19.35.80: 18:55:22.345371
IP 162.222.178.71.35563 > 173.245.58.146.53:
IP 199.71.213.20.40150 > 173.245.58.146.53: 18:55:22.345430
IP 195.175.255.138.62803 > 173.245.58.221.53:
IP 220.213.193.137.52163 > 104.31.188.8.80:
IP 10.40.8.97.8070 > 10.40.8.59.46943:
IP 115.231.91.118.35120 > 173.245.58.146.53:
IP 10.12.11.5.8070 > 10.12.8.106.24514:
```

Takeaways

- You *WILL* null-route
- Linux firewall is awesome
- Sflow for detection

Thanks
and good luck!

(please fill the speaker
excitement form!)

marek@cloudflare.com @majek04